GUESS for Biology

Protein interactions, modules and enrichment, signalling pathways and feedback, expression overlays, with two worked case studies

Contents

1 Why networks? 1

2 Getting the data in 2
2.1 Fetch interactions for a gene list Basic . . . 2
2.2 Open the yeast network and add measurements Basic . . . 3

3 Hubs and bottlenecks 3
3.1 Find the hubs Basic . . . 3
3.2 Find the bottlenecks between modules Intermediate . . . 3

4 Modules and complexes 5
4.1 Find modules and name them Intermediate . . . 5

5 Pathways and regulation 6
5.1 Draw a signalling pathway Basic . . . 6
5.2 Find the feedback loops Advanced . . . 8

6 Measurements on the network 8
6.1 Colour genes by expression Basic . . . 8
6.2 Animate a time course Intermediate . . . 10

7 Case studies 10
7.1 Case study: how pheromone stops the cell cycle Intermediate . . . 10
7.2 Case study: from a hit list to modules Advanced . . . 11

8 Capstone 11
8.1 Capstone: a module browser for a time course Advanced . . . 11

Why networks?

Molecular biology is full of relations: proteins bind, kinases phosphorylate, transcription factors regulate genes, enzymes pass metabolites along, neurons synapse. A network view asks questions that a list of genes can't answer:

- Which proteins hold the system together (hubs and bottlenecks)?

- Which genes work together (modules, complexes)?
- How does a signal travel from receptor to response, and where does it loop back (pathways, feedback)?
- How does a measurement, such as expression after a treatment, map onto that structure?

This book assumes the Kitchen for loading, querying and drawing. Its data are in the cookbooks' `data` folder:

file	what it is
`yeast-modules.gdf`	134 budding-yeast genes from seven well-known complexes and pathways (`module`), with the 1,320 interactions STRING scores at 0.7 or more. Edge `score` is STRING's combined confidence, and `experiments`, `databases`, `textmining` and `coexpression` are the evidence behind it. From string-db.org, CC BY 4.0.
`yeast-go-slim.csv`	Gene Ontology slim terms for those genes (gene, aspect, term), from the Saccharomyces Genome Database, CC BY 4.0.
`pheromone-pathway.csv`	A simplified, directed and signed map of the yeast mating-pheromone pathway, drawn from textbook accounts: who `activates` or `inhibits` whom. For teaching, not a curated resource.
`pheromone-expression-SIMULATED.csv`	**Simulated** log2 fold changes at 15, 30 and 60 minutes after pheromone, for the 134 genes. They are shaped like the known response (mating genes up, G1 cyclins down) but invented: use them to learn the recipes, never as results.

Getting the data in

Fetch interactions for a gene list Basic

You have a list of genes (a screen's hits, a differential-expression list). **You want** the known interactions among them.

Ingredients: an internet connection · `makeFromSTRING`

```
makeFromSTRING("TP53 MDM2 ATM CHEK2 CDKN1A BAX", score=0.7, add=10)
colorize(query)
```

How it works. `makeFromSTRING(genes)` asks the STRING database for interactions among the genes. `species` is an NCBI taxon number (9606 human, the default; 10090 mouse; 4932 yeast; 7227 fly; 6239 worm; 3702 Arabidopsis), and `score` is the minimum combined confidence (0.4 medium, 0.7 high, 0.9 highest). `add=10` also brings in the ten proteins that interact most with your list, marked `query = false`. That's how a screen's hits get connected through genes the screen missed. Every edge records the evidence behind its score: `experiments`, `databases`, `textmining`, `coexpression` and others.

Variations. `makeFromNDEx(uuid)` opens a published network from the NDEx repository (Cytoscape's network commons). Files from Cytoscape open with `makeFromGraphML`, `makeFromGML` or `makeFromCSV` (for SIF-like tables).

Pitfalls. A STRING edge means "associated", not "binds". Text mining and co-expression edges link genes that are studied together or switched on together. For physical interactions, keep the edges with strong experimental evidence, `showOnly(experiments > 0.5)`, or build from a physical-interaction database.

Open the yeast network and add measurements Basic

```
makeFromGDF("yeast-modules.gdf")
makeFromCSV("pheromone-expression-SIMULATED.csv")
FUS1.description, FUS1.lfc60
```

How it works. The second file is a node table (an `Id` column plus values), so GUESS adds its columns to the matching nodes instead of replacing the graph (*Kitchen: Add attributes from a second table*). The same works for any per-gene table: expression, essentiality, abundance, mutation counts. Gene names must match node names. Check the message ("134 updated"), since unmatched IDs become new, unconnected nodes.

Hubs and bottlenecks

Find the hubs Basic

```
makeFromGDF("yeast-modules.gdf")
plotDistrib(degree)
[(n.name, n.module, n.degree) for n in sortBy(degree)[-5:]]
```

How it works. In most molecular networks a few proteins have many partners and most have few. `plotDistrib` shows the shape. The biggest hubs here are proteasome subunits (RPN11 has 45 partners), because a large complex makes every member a partner of many others. The cell-cycle kinase CDC28 is the only hub outside a complex.

Pitfalls. Hubs are partly an artefact of attention: well-studied proteins have more recorded interactions. Compare against `experiments` evidence, or against a less biased data source, before concluding that a hub is biologically central.

Find the bottlenecks between modules Intermediate

You want the proteins that connect otherwise separate parts of the network, which are often signalling or regulatory links.

```
makeFromGDF("yeast-modules.gdf")
colorize(module)
legendFor(module)
sizeBy(betweenness, 6, 40)
forceAtlas2Layout()
interEdges(module).color = gray
[(n.name, n.module, round(n.betweenness)) for n in sortBy(betweenness)[-6:]]
```

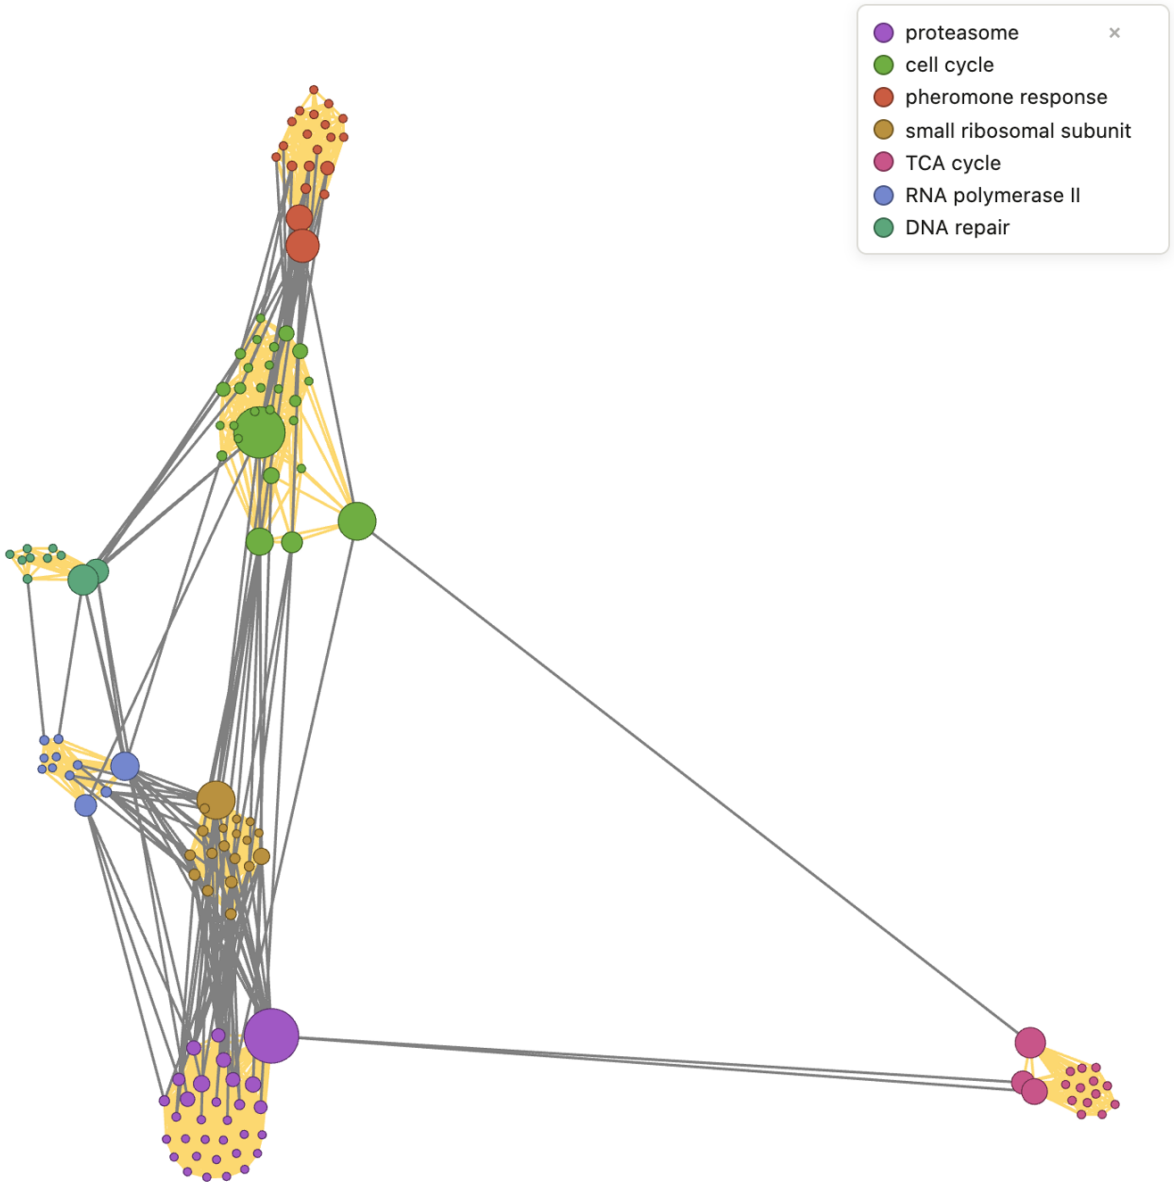


Figure 1: Seven yeast modules, sized by betweenness

How it works. High betweenness means many shortest paths run through a protein. Here the list mixes two kinds of bottleneck. RPN11 (proteasome) and RPS3 (ribosome) are hubs of big complexes that touch many others. CDC28, FAR1 and SKP1 sit where modules meet. FAR1 is the pheromone response's link to the cell cycle: all 11 of its partners outside its own module are cell-cycle proteins.

Variations. `mixingMatrix(module)` counts the edges between each pair of modules. The largest cross-module count here is proteasome–ribosome (48), and 41 of those edges have strong experimental evidence. That isn't necessarily function: very abundant complexes tend to turn up together in pull-down experiments. Check the evidence channels before reading a cross-module edge as

cross-talk.

Modules and complexes

Find modules and name them Intermediate

You want groups of genes that work together, and to know what each group does.

Ingredients: `yeast-modules.gdf`, `yeast-go-slim.csv` · `leidenClusters` · `readSets` · `enrichment`

```
makeFromGDF("yeast-modules.gdf")
leidenClusters()
go = readSets("yeast-go-slim.csv", item="gene", set="term",
              where={"aspect": "process"})
r = enrichment(leiden_class == 0, go, top=5)
```

How it works. `leidenClusters()` finds communities (*Kitchen: Communities*). `readSets` reads a table with one row per gene and term into a dictionary of gene sets. `where=` keeps only biological-process terms. `enrichment(nodes, sets)` asks, for each term, whether the community has more of its genes than chance. For each term it reports:

column	meaning
`k/n`	community genes with the term / annotated community genes
`K/N`	genes with the term / annotated genes overall
`fold`	(k/n) ÷ (K/N): how over-represented
`p`	the chance of so many or more by chance (hypergeometric)
`q`	the false discovery rate over all terms tested (Benjamini–Hochberg)

The largest community has 28 of its 31 genes in "protein catabolic process" (q ≈ 2e-15): it's the proteasome. `r` holds every row, and `r[0]["members"]` is a set of the genes, so `r[0]["members"].color = red` shows them.

Variations. Name every community at once:

```
makeFromGDF("yeast-modules.gdf")
leidenClusters()
go = readSets("yeast-go-slim.csv", item="gene", set="term",
              where={"aspect": "process"})
for grp in groupBy(leiden_class):
    best = enrichment(grp, go, show=False)[0]
    print(grp[0].leiden_class, len(grp), best["set"], "%.0e" % best["q"])
```

Pitfalls.

- **Choose the background deliberately.** By default it's every node that has at least one annotation. Against the whole genome, everything in a hand-picked network looks enriched. Pass `background=` to choose.

- **These data are unrealistically clean.** The genes were picked from known complexes, so communities match them exactly (`compareClusters(leiden_class, module)` is 1.0). Real gene lists give messier modules, and that's when enrichment earns its keep.

Pathways and regulation

Draw a signalling pathway Basic

You have a directed, signed pathway. **You want** a readable picture, from receptor to response.

Ingredients: `pheromone-pathway.csv` · `dotLayout` · edge queries

```
makeFromCSV("pheromone-pathway.csv")
labelBy(name)
g.nodes.labelpos = "inside"
fitNodesToLabels()
dotLayout()
(effect == "inhibits").color = red
(effect == "activates").color = darkgray
STE2.color = orange
STE12.color = purple
```

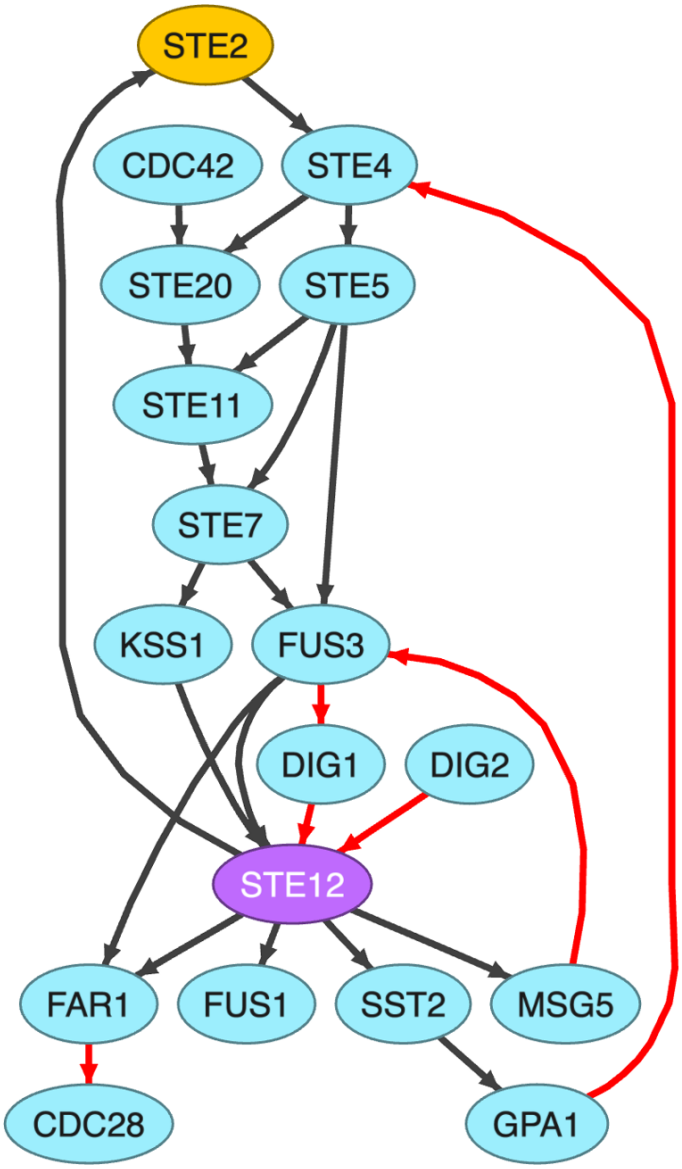


Figure 2: The pheromone pathway, laid out by Graphviz

How it works. `dotLayout()` (Graphviz) ranks nodes top to bottom by following the arrows. Feedback edges point back up, which makes loops easy to spot. `fitNodesToLabels()` sizes each node to hold its name. Do it *before* the layout, so the layout leaves room for the bigger nodes.

Variations. `descendants(STE2)` is everything downstream of the receptor. `ancestors(CDC28)` is everything that can influence the cell-cycle kinase. `shortestPath(STE2, CDC28)` is the most direct route.

Find the feedback loops Advanced

You want the cycles in a pathway, and whether each one amplifies the signal (positive feedback) or damps it (negative feedback).

```
makeFromCSV("pheromone-pathway.csv")
import networkx as nx
G = toNetworkX()
for cyc in sorted(nx.simple_cycles(G), key=len)[:6]:
    steps = list(zip(cyc, cyc[1:] + cyc[:1]))
    signs = [G.edges[a, b]["effect"] for a, b in steps]
    kind = "negative" if signs.count("inhibits") % 2 else "positive"
    print(kind, " -> ".join(cyc))
```

How it works. `toNetworkX()` hands the graph to networkx, whose `simple_cycles` lists every cycle. A loop's sign is the product of its edges' signs, so an odd number of inhibitions makes it negative. The shortest loop, FUS3 → STE12 → MSG5 -| FUS3, is negative: the response switches on a phosphatase that turns the kinase off. The loop through the receptor (STE12 → STE2) is positive: responding cells make more receptor.

Pitfalls. Real pathways have many overlapping cycles (24 in this small map). List the shortest, and look at which edges they share: `feedbackEdges()` gives the fewest edges whose removal breaks every cycle.

Measurements on the network

Colour genes by expression Basic

```
makeFromGDF("yeast-modules.gdf")
makeFromCSV("pheromone-expression-SIMULATED.csv")
forceAtlas2Layout()
colorBy(lfc60, blue, red, center=0)
legendFor(lfc60)
```

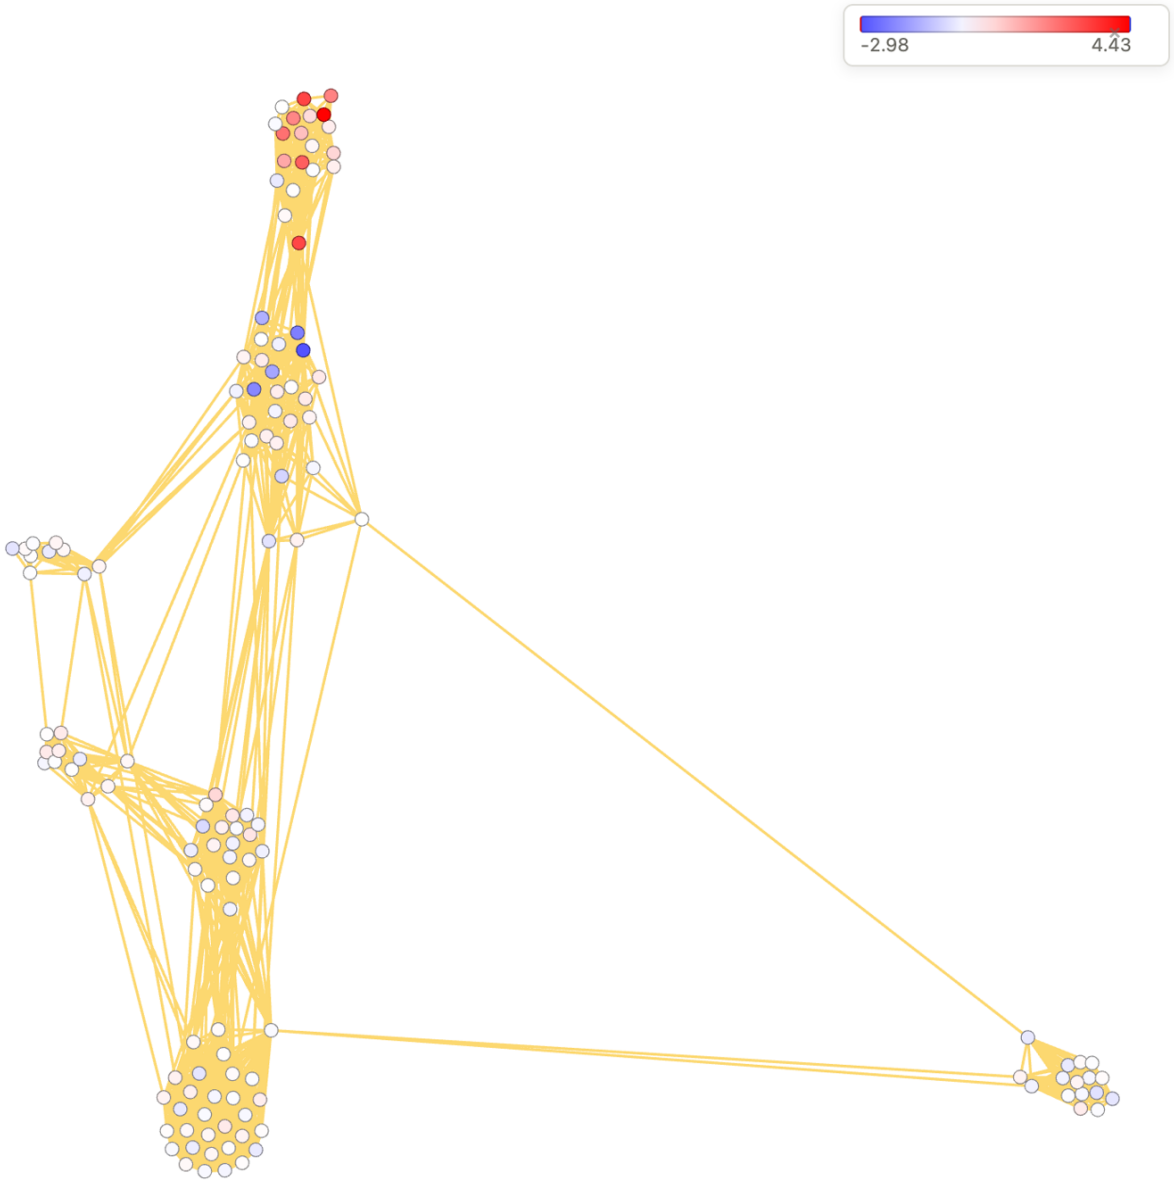


Figure 3: Simulated expression at 60 minutes, blue = down, red = up

How it works. `colorBy(field, low, high)` maps numbers to a gradient. `center=0` makes it *diverging*: zero is white, and the two ends are equally far from zero, so a gene down 2 is as deep a blue as a gene up 2 is red. That's the right choice for fold changes, where zero means "no change" and both directions matter.

Pitfalls. Without `center=`, a three-colour gradient puts its middle colour halfway between the lowest and highest values. Here the values run from -3.0 to 4.4, so "white" would fall at $+0.7$ and every unchanged gene would be tinted. Always pin the centre for signed data.

Animate a time course Intermediate

```
makeFromGDF("yeast-modules.gdf")
makeFromCSV("pheromone-expression-SIMULATED.csv")
forceAtlas2Layout()
for t in (lfc15, lfc30, lfc60):
    colorBy(t, blue, red, center=0)
    ss(t.name)
ls("lfc15")
morph("lfc30", 1500)
```

How it works. Each time point is saved as a state with its own colours, and `morph` animates between them, so the response is seen spreading through the network. `addSlide` instead of `ss` makes a presentation. To show all time points at once, `pieChart([lfc15, lfc30, lfc60])` draws each gene as a pie of its values. Pies work best with positive amounts (abundances, counts), not signed changes.

Case studies

Case study: how pheromone stops the cell cycle Intermediate

A haploid yeast cell that senses mating pheromone stops dividing in G1, so it can fuse with a partner. Starting from the network alone, how does the signal reach the cell-cycle machinery?

```
makeFromGDF("yeast-modules.gdf")
colorize(module)
p = shortestPath(STE2, CDC28)
[n.name for n in p if n in g.nodes]
```

In the STRING network the shortest route from the receptor STE2 to the cyclin-dependent kinase CDC28 runs through BEM1, a polarity scaffold shared by mating and budding. Shortest paths in an association network are suggestive, not mechanistic: an association can be any kind of link.

```
showOnly((module == "pheromone response") | (module == "cell cycle"))
forceAtlas2Layout()
sizeBy(betweenness, 6, 40)
bridging = interEdges(module)
bridging.color = red
ends = [n for e in bridging for n in (e.node1, e.node2)]
sorted(set(n.name for n in ends if n.module == "pheromone response"))
```

The 28 edges between the two modules touch ten pheromone genes, but FAR1 alone accounts for 11 of them, all to cell-cycle proteins. That points to the known mechanism (FUS3 activates FAR1, which inhibits CDC28 with its G1 cyclins), which the directed pathway map makes explicit:

```
makeFromCSV("pheromone-pathway.csv")
[n.name for n in shortestPath(STE2, CDC28) if n in g.nodes]
```

Following the arrows, the route is STE2 → STE4 → STE5 → FUS3 → FAR1 -| CDC28: receptor, G-protein, scaffold, MAP kinase, then the CDK inhibitor.

Case study: from a hit list to modules Advanced

A screen gives you genes. With a network connection, the full workflow is:

```
hits = "RPN11 RPT5 PRE1 CDC28 CLN2 FAR1 FUS3 STE12 RPO21 MSH2"
makeFromSTRING(hits, species=4932, add=20)
leidenClusters()
```

The bundled network stands in for that result here:

```
makeFromGDF("yeast-modules.gdf")
subgraph(experiments > 0.5, "physical")
leidenClusters()
go = readSets("yeast-go-slim.csv", item="gene", set="term",
              where={"aspect": "process"})
for grp in groupBy(leiden_class)[:5]:
    best = enrichment(grp, go, show=False)[0]
    print(len(grp), best["set"], "q=%.0e" % best["q"])
```

`subgraph(experiments > 0.5, "physical")` opens a graph of just the edges with strong experimental evidence (1,028 of 1,320) and the genes they join (`use("main")` goes back). Communities are then re-found and named by their top term. Which modules survive with physical evidence alone, and which were held together by co-expression and text mining? Ask that of any network built from mixed evidence. Here all seven survive, each still enriched for the right process. With a real hit list, expect some modules to dissolve.

Pitfall. A subgraph of edges loses any gene left with no strong edge. `subgraph(experiments > 0.5)` keeps 134 genes here, but on noisier data count the nodes before and after, and treat those that drop out as “no physical evidence”, not “no interaction”.

Capstone

Capstone: a module browser for a time course Advanced

The problem. A lab has a time course after a treatment and wants to browse it through the interaction network. Which modules respond? When? Which genes lead? This recipe builds a module browser:

- modules are found and named by their top GO term, and a drop-down picks one;
- a slider steps through the time points, recolouring every gene;
- a checkbox switches between all evidence and experimental evidence only;
- a click on a gene shows its description and its values at every time point;
- a legend shows the diverging scale.

```
makeFromGDF("yeast-modules.gdf")
makeFromCSV("pheromone-expression-SIMULATED.csv")
forceAtlas2Layout()
_ = leidenClusters()
go = readSets("yeast-go-slim.csv", item="gene", set="term", where={"aspect":
 ↪  "process"})
names = {}
```

```
for grp in groupBy(leiden_class):
    names[enrichment(grp, go, show=False)[0]["set"]] = grp[0].leiden_class
times = [lfc15, lfc30, lfc60]
state = {"t": 2}

def recolour(i):
    state["t"] = int(i)
    colorBy(times[state["t"]], blue, red, center=0)
    legendFor(times[state["t"]], title=f"log2 change at {[15, 30,
 ↪  60][state['t']]} min")

def pick_module(term):
    members = (leiden_class == names[term])
    g.nodes.labelvisible = false
    members.labelvisible = true
    center(members)
    up = len(members & (times[state["t"]] > 1))
    info.text = f"{term}: {len(members)} genes, {up} up more than 2-fold"

def evidence(only):
    if only:
        showOnly(experiments > 0.5)
    else:
        showAll()

def gene(n):
    values = ", ".join(f"{t}: {getattr(n, f.name):+.1f}" for t, f in zip((15, 30,
 ↪  60), times))
    info.text = f"{n.name} ({n.module}): {n.description}; {values}"

info = Label("pick a module, or click a gene")
ui.dock(Panel("Modules", Dropdown(sorted(names), label="module",
 ↪  onChange=pick_module),
              Slider(0, 2, 2, label="time point (15, 30, 60 min)",
 ↪  onChange=recolour),
              Checkbox("experimental evidence only", onChange=evidence), info,
 ↪  position="bottom"))
graphevents.clickNode = gene
recolour(2)
center()
```

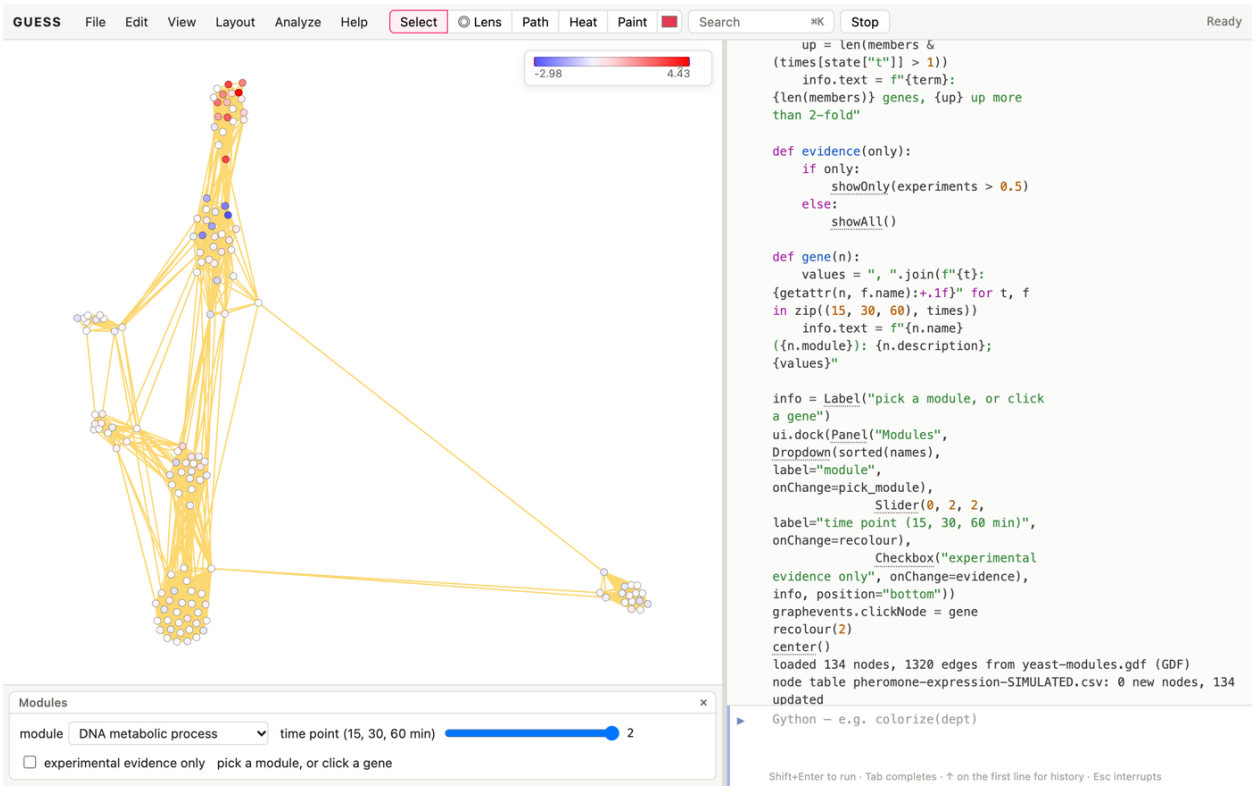


Figure 4: A module browser: simulated expression at 60 minutes on the yeast network

How it works.

- **Naming the modules.** The modules come from `leidenClusters()`, and each is named by its top enriched term (the module-naming recipe), so the drop-down lists biology (“signal transduction”, “protein catabolic process”), not numbers.
- **The slider.** It indexes the three expression fields. `colorBy(..., center=0)` and `legendFor` redraw the diverging scale for each time point.
- **The checkbox.** It toggles `showOnly(experiments > 0.5)`, the physical-evidence view from the second case study.
- **The click.** It reads a gene’s fields into the label.

As before, the expression values are **simulated**: substitute your own table, with an `Id` column of gene names, in the second line.