GUESS for Epidemiology

Contact networks, simulated outbreaks and the epidemic threshold, vaccination and class bubbles, transmission trees and superspreading, backward tracing, effective distance through air travel, with an outbreak control room

Contents

1 Why networks? 1

2 Contact networks 2
2.1 Load a school's contact network Basic . . . 2
2.2 Your contacts have more contacts than you Basic . . . 4

3 Simulating outbreaks 4
3.1 Simulate an outbreak Basic . . . 4
3.2 One run is an anecdote Intermediate . . . 6
3.3 The epidemic threshold Intermediate . . . 6
3.4 Where an outbreak starts matters Intermediate . . . 7

4 Interventions 7
4.1 Whom to vaccinate first? Intermediate . . . 7
4.2 Do class bubbles work? Intermediate . . . 8

5 Case data 9
5.1 A transmission tree from a line list Basic . . . 9
5.2 Superspreading: R and k Intermediate . . . 11

6 Travel 11
6.1 How far is it, for an epidemic? Advanced . . . 11

7 Case study 14
7.1 Case study: trace backwards Advanced . . . 14

8 Capstone 14
8.1 Capstone: an outbreak control room Advanced . . . 14

Why networks?

Classic epidemic models assume everyone mixes with everyone: each infected person is equally likely to meet each susceptible one. Real infections travel along contacts: a classroom, a household,

a flight. The network view asks questions those models can't:

- Who is most at risk, and who would spread an infection furthest?
- How contagious must a disease be to cause a large outbreak *in this population*?
- Whom should we vaccinate first, and does keeping classes apart work?
- In an outbreak that has happened, who infected whom, and how uneven was it?
- Which cities will an epidemic reach next?

This book assumes the Kitchen. Its datasets are **simulated**, so they can be shared freely and the patterns in them are known:

- `school-contacts-SIMULATED.gdf`: a primary school of 246 pupils in ten classes (1A to 5B) and their ten teachers. Edge `weight` is minutes of close contact in a day, as studies with wearable sensors record. Its structure follows what those studies report: most contact within a class, less within a grade, little between grades.
- `outbreak-cases-SIMULATED.csv` and `outbreak-transmissions-SIMULATED.csv`: a line list of a 64-day respiratory outbreak in a town. Gatherings were banned on day 28.

The travel recipe uses the real world airline network from *GUESS for Infrastructure*.

In the simulations, one step is one day. With `gamma=0.2`, an infected person stays infectious for five days on average.

Contact networks

Load a school's contact network Basic

You want to see who meets whom at school, and how much.

Ingredients: `school-contacts-SIMULATED.gdf` · `colorize` · `sizeBy`

```
makeFromGDF("school-contacts-SIMULATED.gdf")
colorize(room)
sizeBy(degree, 4, 20)
(role == "teacher").style = 1
legendFor(room)
[(grp[0].room, len(grp)) for grp in groupBy(room)][:3]
```

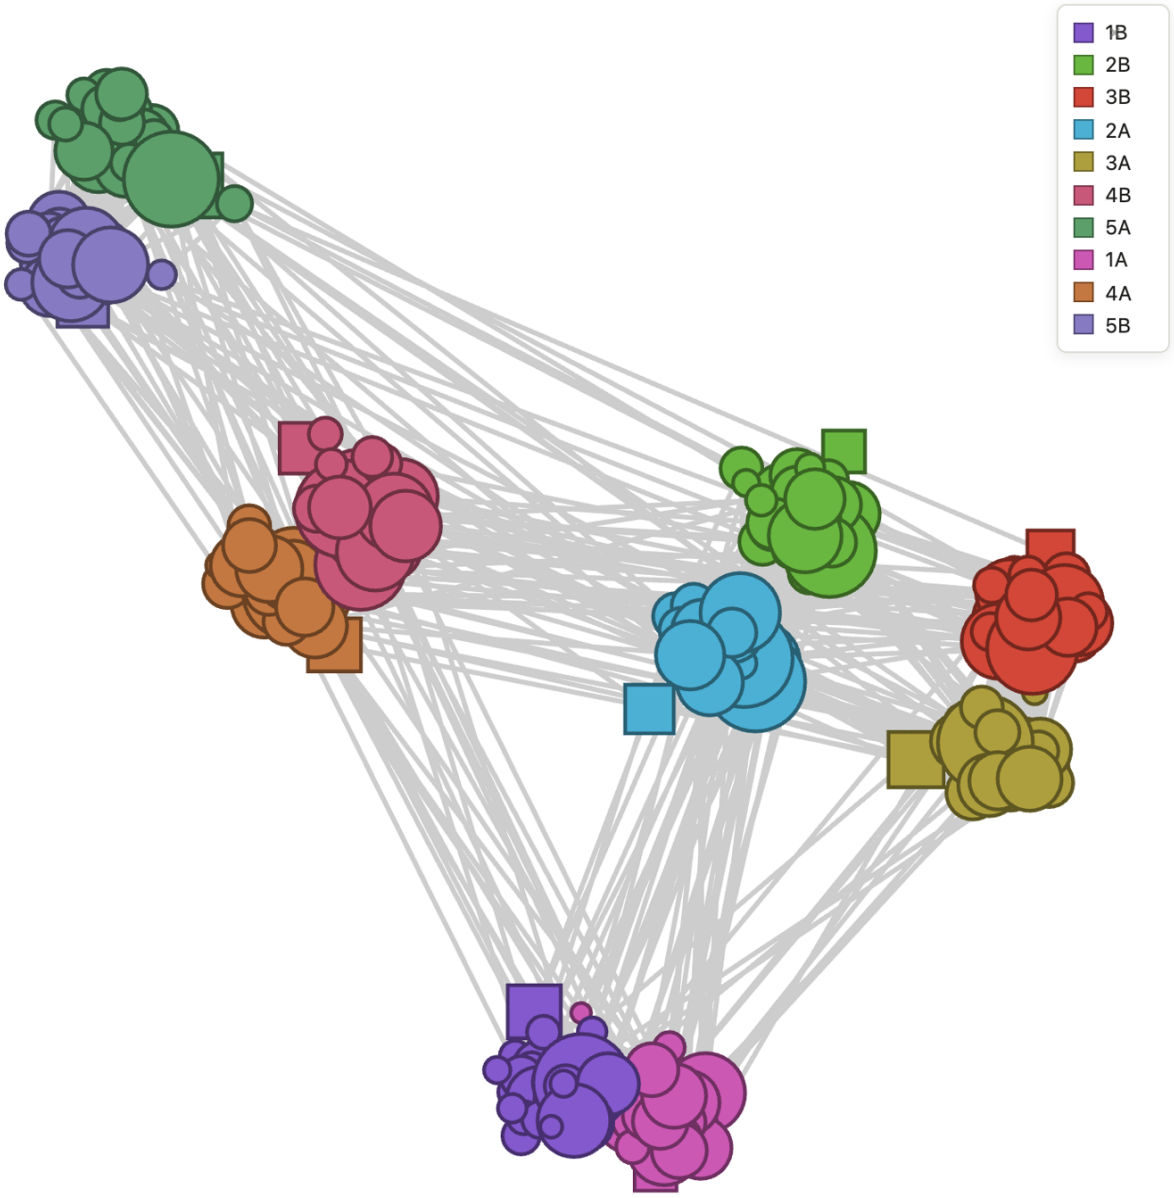


Figure 1: A school's contact network: one colour per class, teachers as squares

How it works. Each node is a person, and each edge a pair who were in close contact during the day. `room` is the class (`1A` to `5B`), `grade` the year, and `role` pupil or teacher. The classes are plain to see: 1,633 of the 2,125 contacts, and 91% of the contact minutes, are within a class. The rest, between classes, is what lets an infection leave one. A person has 16.6 contacts on average, from 2 to 38.

Pitfalls. A contact network is a sample of a day. Sensors miss short or distant contacts, and a different day gives a different network. Studies that recorded several days find the strong, within-class ties repeat and the weak ones change.

Your contacts have more contacts than you Basic

```
makeFromGDF("school-contacts-SIMULATED.gdf")
import statistics
mine = statistics.mean(n.degree for n in g.nodes)
theirs = statistics.mean(statistics.mean(m.degree for m in n.getNeighbors()) for
 ↪  n in g.nodes)
round(mine, 1), round(theirs, 1)
```

How it works. The average person here has 16.6 contacts, but the average contact of a person has 19.9. That's the *friendship paradox* (Feld 1991): well-connected people appear in many contact lists, so a contact chosen at random is likely to be one of them. It holds for 68% of the people in this school.

It matters for epidemics, because an infection reaches people through their contacts too. The people an outbreak reaches early are better connected than average, and the paradox gives two practical tricks used below: vaccinating a random person's *contact* (acquaintance immunisation) and tracing a case's *infector* (backward tracing).

Simulating outbreaks

Simulate an outbreak Basic

You want to watch an infection spread through the school from one pupil.

Ingredients: a contact network · `spread` · `showTime`

```
makeFromGDF("school-contacts-SIMULATED.gdf")
counts = spread([p001], "SIR", beta=0.03, gamma=0.2, seed=1)
colorBy(infected_at, yellow, red)
(infected_at < 0).color = lightgray
center()
```

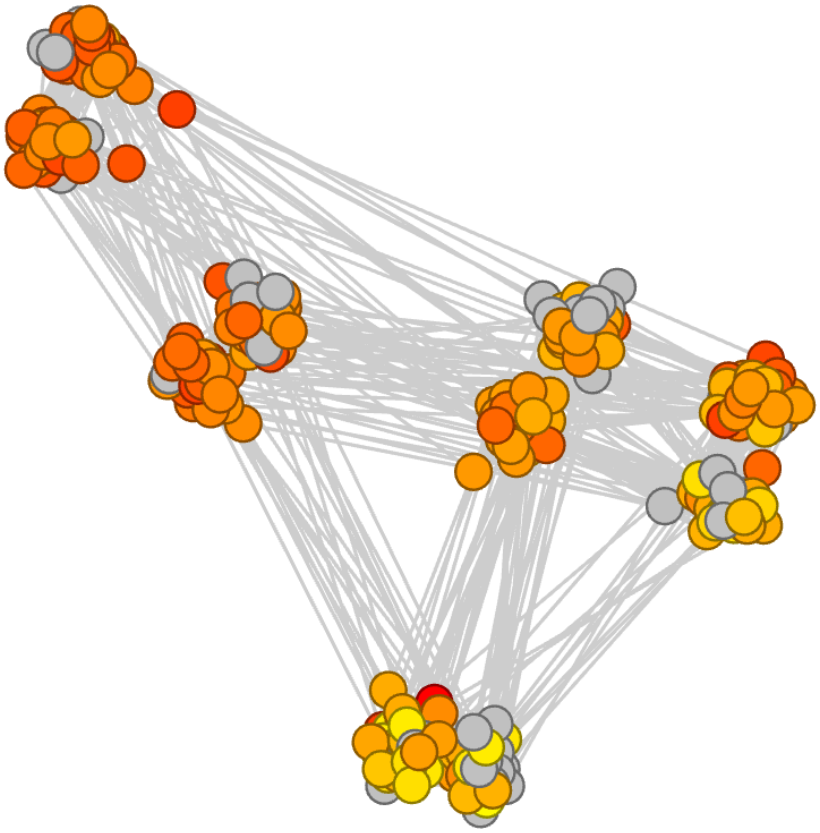


Figure 2: One simulated outbreak, coloured by the day each person was infected (yellow early, red late, grey never)

How it works. `spread` runs an SIR model: each day, each infected person infects each susceptible contact with probability `beta`, and recovers with probability `gamma`. It stores the day each person was infected (`infected_at`, -1 for never) and the final `state`, and plots the susceptible, infected and recovered counts in the chart panel. This run started in class 1A and reached 207 of the 256 people, in every class, over 49 days. 46 people were infectious at the peak, on day 20. With `seed=0` it fizzles out at once: chance matters most at the start.

Variations. `showTime(10, infected_at)` shows only the people infected by day 10; drag a slider over it to replay the outbreak. `spread(..., weight=weight, beta=0.002)` makes the risk depend on contact minutes: an edge of 20 minutes then transmits with probability $1 - 0.998^{20}$, about 4% a day.

One run is an anecdote Intermediate

You want the distribution of outcomes, not one story, and to know who is most at risk.

Ingredients: a contact network · `outbreaks`

```
makeFromGDF("school-contacts-SIMULATED.gdf")
sizes = outbreaks(300, beta=0.03, gamma=0.2)
sizes.mean, sizes.large, sizes.meanLarge
colorBy(risk, white, red)
```

How it works. `outbreaks(runs)` repeats the simulation from a random pupil or teacher each time and charts the final sizes. The chart has two humps. 44% of the outbreaks fizzle out after a few cases. The other 56% take off and reach 204 people on average, four in five of the school. The mean, 116, describes almost no actual outbreak, which is why the result also reports `.large`, the share that take off (reaching more than 10% of the nodes), and `.meanLarge`.

`outbreaks` also stores each person's `risk`: the share of runs that reached them. Risk follows contacts closely here (correlation 0.86 with `degree`): the best-connected pupils were infected in 56% of the runs, the least connected in 12% to 19%.

Pitfalls. The model knows nothing about immunity, age or symptoms, and real outbreaks are seeded by someone who caught it elsewhere, not by a random pupil. Use simulations to compare options on the same network, not to forecast case counts.

The epidemic threshold Intermediate

You want to know how contagious an infection must be to cause a large outbreak in this school.

```
makeFromGDF("school-contacts-SIMULATED.gdf")
import statistics
k = statistics.mean(n.degree for n in g.nodes)
k2 = statistics.mean(n.degree ** 2 for n in g.nodes)
for b in [0.01, 0.02, 0.03, 0.05]:
    T = b / (b + 0.2 - b * 0.2)
    sz = outbreaks(200, beta=b, gamma=0.2, show=False)
    print(b, round(T * (k2 - k) / k, 2), sz.large)
```

How it works. The *basic reproduction number* R_0 is the number of people a typical case infects in a fully susceptible population. Above 1, an outbreak can take off. On a network it depends on two

things:

- the *transmissibility* T, the chance an infected person passes it to a given contact before recovering: $T = \beta / (\beta + \gamma - \beta\gamma)$ in this model;
- the spread of contacts: $R_0 = T\,(\langle k^2 \rangle - \langle k \rangle) / \langle k \rangle$, where $\langle k \rangle$ is the mean number of contacts and $\langle k^2 \rangle$ the mean of its square.

beta	R_0	share of outbreaks that take off
0.01	0.90	4%
0.02	1.73	45%
0.03	2.50	54%
0.05	3.89	77%

Below the threshold, large outbreaks are rare. Above it, they become likely, but never certain: an outbreak can still die out by chance in its first few cases.

Why $\langle k^2 \rangle$? Someone infected was reached through a contact, so is likely well connected (the friendship paradox again), and passes it on to more people than the average person would. The more uneven the contacts, the bigger $\langle k^2 \rangle$ and the lower the threshold. In networks with very uneven contacts, like sexual contact networks, the threshold can be close to zero (Pastor-Satorras and Vespignani 2001).

Where an outbreak starts matters Intermediate

```
makeFromGDF("school-contacts-SIMULATED.gdf")
for n in sortBy(degree)[-2:] + sortBy(degree)[:2]:
    print(n.name, n.degree, outbreaks(100, seeds=[n], beta=0.03, gamma=0.2,
     ↪  show=False).large)
```

How it works. `seeds=[n]` starts every run from the same person. Started from one of the best-connected pupils (37 or 38 contacts), 81% to 90% of outbreaks take off. Started from the least connected (2 or 3 contacts), 14% to 20% do. A well-connected first case gets many chances before recovering.

Variations. Seeding from each person in turn and storing the share (`n.reach = ...`) maps who would spread an infection furthest. Kitsak et al. (2010) found that `coreness` (the k-core a node is in) predicts it better than degree in many networks. Try both.

Interventions

Whom to vaccinate first? Intermediate

You have enough vaccine for 20% of the school. **You want** to know whom to give it to.

Ingredients: a contact network · `outbreaks(immune=...)`

```
makeFromGDF("school-contacts-SIMULATED.gdf")
import random
rng = random.Random(1)
```

```
dose = round(0.2 * len(g.nodes))
chosen = {"random": rng.sample(list(g.nodes), dose), "most contacts":
 ↪  sortBy(degree)[-dose:]}
acq = set()
while len(acq) < dose:
    acq.add(rng.choice(list(rng.choice(list(g.nodes)).getNeighbors())))
chosen["acquaintance"] = list(acq)
for how, who in chosen.items():
    print(how, outbreaks(300, beta=0.03, gamma=0.2, immune=who, show=False))
```

How it works. `immune=` marks people who can't be infected. With no vaccine, outbreaks reach 116 people on average:

who gets the 51 doses	mean outbreak	share that take off
nobody	116	56%
51 people at random	62	48%
a random contact of 51 random people	42	42%
the 51 with the most contacts	12	14%

Targeting the best-connected people is by far the best, but it needs the whole contact network, which is rarely known. *Acquaintance immunisation* (Cohen, Havlin and ben-Avraham 2003) needs none: pick people at random and vaccinate one of their contacts. The friendship paradox does the targeting. Those chosen here have 20.4 contacts on average, against 16.3 for random choice.

Pitfalls. "Most contacts" is measured on one day's network. Targets chosen on one day protect less on another, when the weak ties have changed. And vaccines aren't perfect: remove a random share of `immune` to model that.

Do class bubbles work? Intermediate

You want to know what keeping classes apart (no mixing at breaks and lunch) would do.

```
makeFromGDF("school-contacts-SIMULATED.gdf")
ss("school")
cross = [e for e in g.edges if e.node1.room != e.node2.room]
len(cross), len(g.edges)
remove(cross)
outbreaks(300, beta=0.03, gamma=0.2)
ls("school")
import random
remove(random.Random(2).sample(list(g.edges), len(cross)))
outbreaks(300, beta=0.03, gamma=0.2)
ls("school")
```

How it works. The 492 contacts between classes are 23% of all contacts. Removing them leaves ten separate classes, and the mean outbreak falls from 116 to 10: an outbreak stays in the class it starts in. Removing the same number of contacts at random barely helps (72). What matters is not how many contacts are cut but *which*: the between-class contacts are the bridges. `ss` and `ls` save and restore the school between experiments.

Variations. Cut only the contacts between grades (`e.node1.grade != e.node2.grade`) and see how much of the benefit remains, or keep teachers' contacts across classes.

Case data

A transmission tree from a line list Basic

You have a *line list*, one row per case, and the results of contact tracing: who infected whom, and where. **You want** the outbreak as a network.

Ingredients: `outbreak-cases-SIMULATED.csv` (Id, onset, age, ward) · `outbreak-transmissions-SIMULATED.csv` (Source, Target, setting) · `makeFromCSV`

```
makeFromCSV("outbreak-transmissions-SIMULATED.csv",
 ↪  "outbreak-cases-SIMULATED.csv", directed=true)
treeLayout(horizontal=true, spacing=6)
for n in g.nodes: n.x = 20 * n.onset
colorize(setting)
colorBy(onset, yellow, red)
sizeBy(outdegree, 4, 24)
legendFor(setting)
center()
```

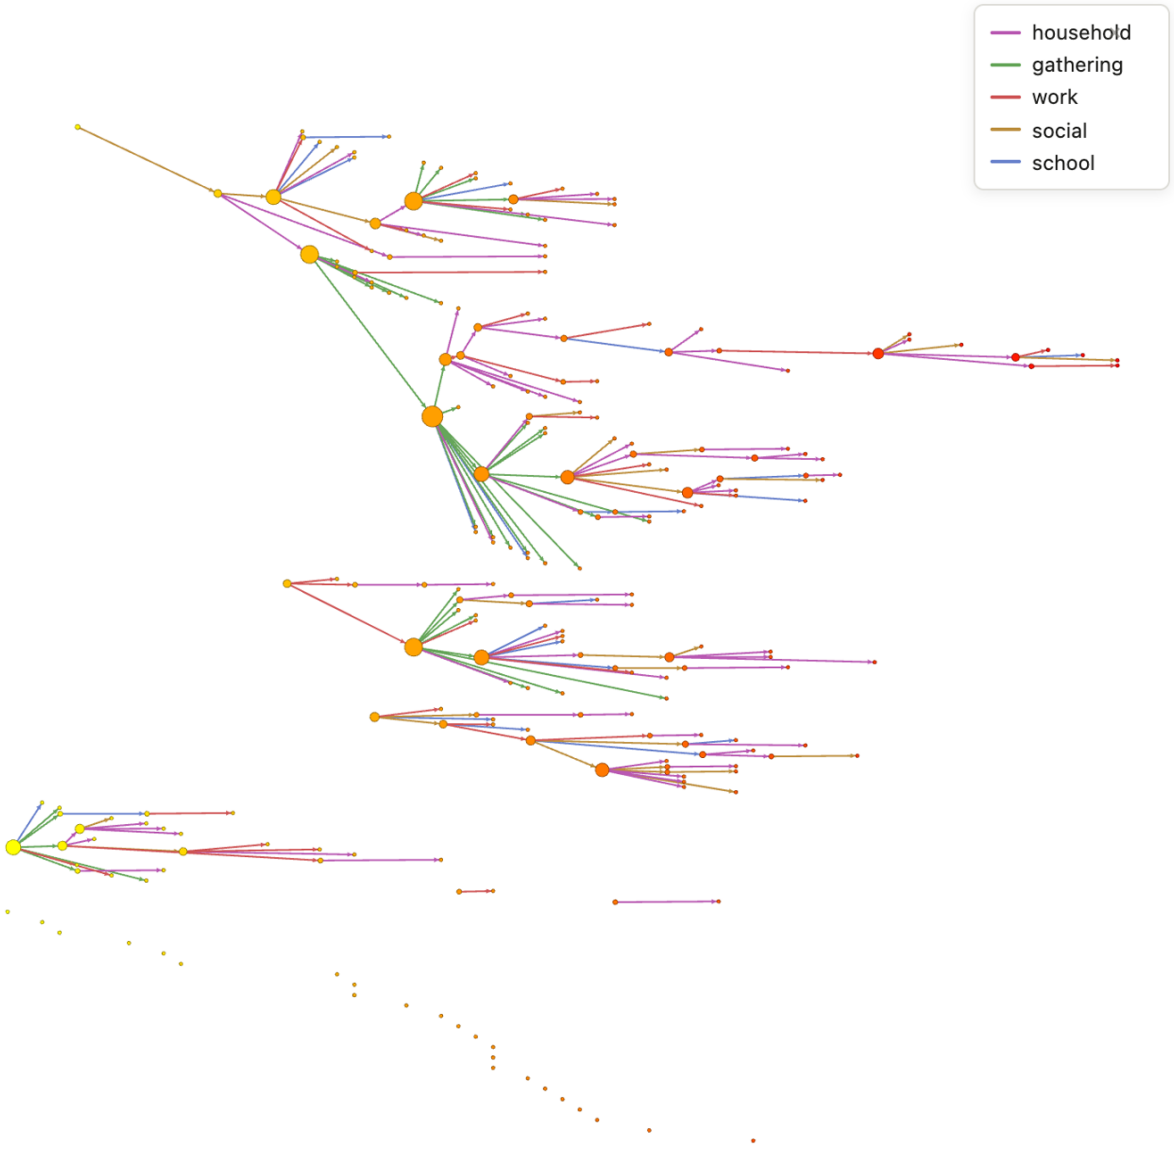


Figure 3: The transmission tree: edges coloured by setting, cases by onset day, sized by how many they infected

How it works. The transmissions table is an edge list (Source infected Target), and the cases table adds each case's fields, including cases with no known infector. Of 234 cases, 29 are roots: the first three, and those contact tracing couldn't link. `treeLayout(horizontal=true)` draws each chain of transmission as a tree growing to the right, and the loop then moves each case to its onset day, so the drawing is also a timeline.

```
plotHistogram(onset, 13)
```

The histogram of `onset` is the *epidemic curve*: cases rise to a peak around day 30 and fall after the

ban on gatherings on day 28, with a lag of about one serial interval.

A *serial interval*, the days from an infector's symptoms to their case's, is one line per edge:

```
addEdgeField("interval", INTEGER, 0)
for e in g.edges: e.interval = e.node2.onset - e.node1.onset
plotHistogram(interval)
```

They average 5.0 days here (1 to 15), in every setting.

Pitfalls. Contact tracing mixes evidence and guesswork. When two possible infectors are in the same household, the one recorded may be wrong. Treat a transmission tree as a hypothesis, and check it against onset dates: an infector whose onset is after their case's is a red flag.

Superspreading: R and k Intermediate

You want to know how many people each case infected, and how uneven that was.

Ingredients: a transmission tree · `offspring`

```
makeFromCSV("outbreak-transmissions-SIMULATED.csv",
 ↪  "outbreak-cases-SIMULATED.csv", directed=true)
offspring(onset < 28)
offspring((onset >= 28) & (onset < 55))
```

How it works. A case's *offspring* is how many it infected, its outdegree in the tree. `offspring()` reports their mean, the *reproduction number* R, and fits a negative binomial distribution to them. Its *dispersion* k measures unevenness: large k means everyone infects about R people, and k below 1 means a few *superspreaders* cause most infections. SARS had k ≈ 0.16 (Lloyd-Smith et al. 2005), COVID-19 around 0.1 to 0.5.

cases with onset	R	k	infected no one	caused by the top 20%
before day 28	1.39	0.21	66%	89%
days 28–54	0.54	0.24	75%	92%

Before the ban on gatherings, each case infected 1.39 people on average: the outbreak was growing. After, 0.54: it was shrinking. Throughout, most cases infected no one and a fifth caused nine in ten infections. `offspring` also stores the counts in the `offspring` field and charts them.

Pitfalls.

- *Censoring.* Cases near the end of the data haven't had time to infect anyone yet, which is why the recipe stops at day 54.
- *Missed links.* Cases whose infector wasn't found (29 here) lower R: an infection that tracing missed counts for no one.

Travel

How far is it, for an epidemic? Advanced

You want to know which cities a new disease will reach first after it emerges in Mexico City, as H1N1 influenza did in 2009.

Ingredients: `world-airports.gdf` (edge `weight` = airline routes) · `effectiveDistance`

```
makeFromGDF("world-airports.gdf")
effectiveDistance(MEX)
for c in [GDL, CUN, LAX, JFK, MAD, LHR, NRT, SYD]: print(c.name, c.city,
 ↪  c.eff_distance)
radialLayout(MEX, eff_tree == true)
import math
for n in eff_distance != None:
    a = math.atan2(n.y, n.x)
    n.x, n.y = 60 * n.eff_distance * math.cos(a), 60 * n.eff_distance *
 ↪  math.sin(a)
showOnly(eff_tree == true)
colorBy(eff_distance, red, yellow)
center()
```

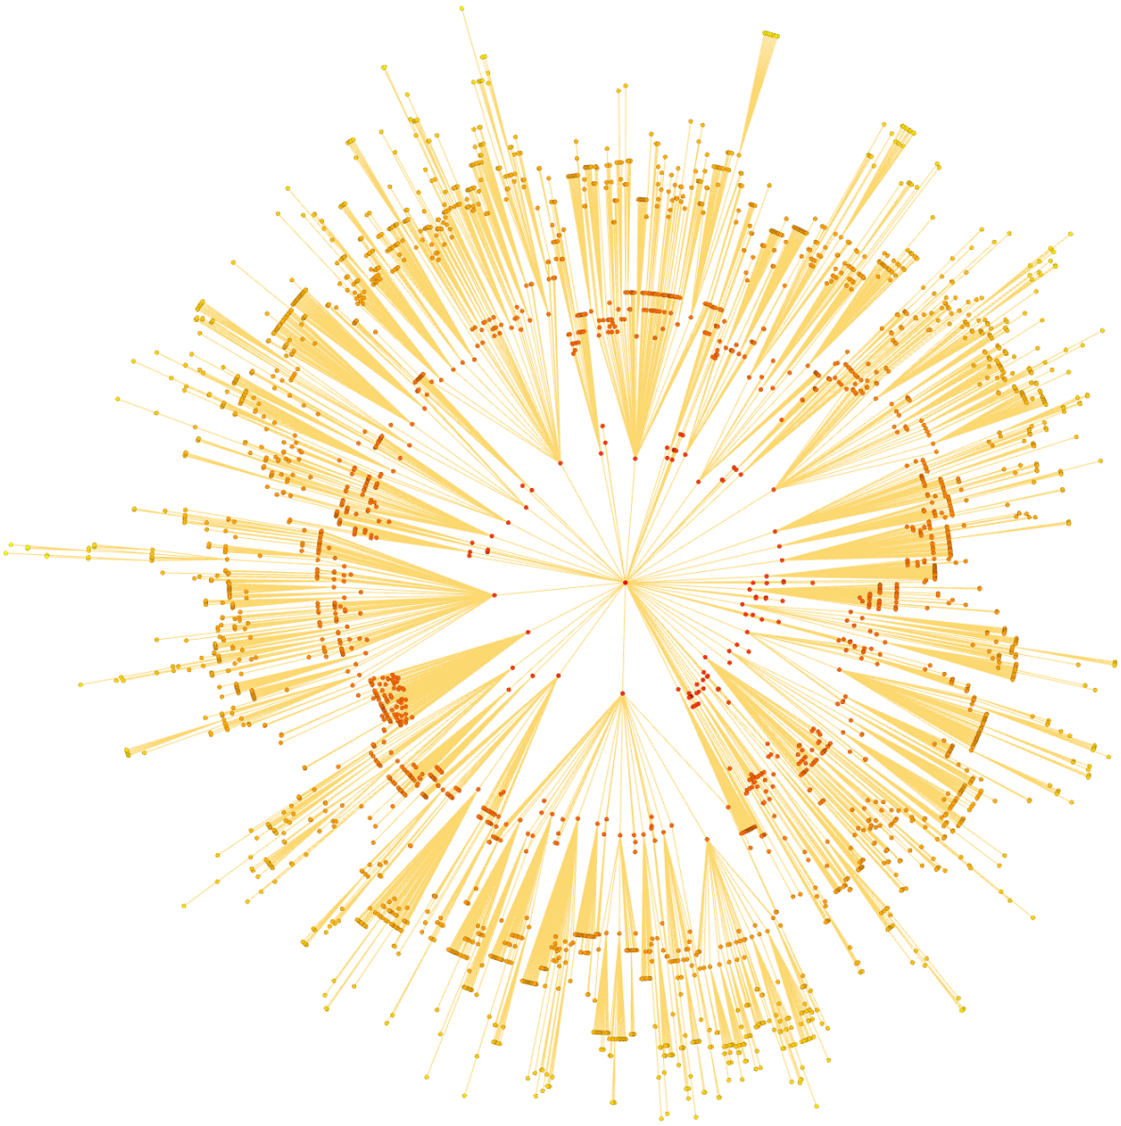


Figure 4: The world from Mexico City: each airport at its effective distance, along the paths an epidemic would most likely take

How it works. Brockmann and Helbing (2013) noticed that the dates on which SARS and H1N1 reached each country fit poorly to distance in kilometres, but almost perfectly to an *effective distance* through the air network. From an airport m, if a share P of its traffic goes to n, the step m $\to$ n has length $1 - \ln P$. Busy routes are short, and a route among hundreds is long. `effectiveDistance` finds the shortest total from the origin and marks the tree of shortest routes in `eff_tree`.

airport	km from Mexico City	effective distance
Guadalajara (GDL)	459	5.10
Cancún (CUN)	1,284	4.88
Los Angeles (LAX)	2,500	4.55
New York (JFK)	3,365	5.10
Madrid (MAD)	9,067	5.10
London (LHR)	8,904	5.57
Tokyo (NRT)	11,254	7.18
Sydney (SYD)	12,983	9.95

To an epidemic, Madrid is as close to Mexico City as Guadalajara, and Los Angeles closer than either. The drawing puts each airport at its effective distance from the centre: the world reorganises into rings, and arrival dates follow the rings.

Pitfalls. The weights here count airline route listings, not passengers, so they're only a stand-in for traffic. Brockmann and Helbing used passenger flows. With passenger numbers in `weight`, the same command gives their measure.

Case study

Case study: trace backwards Advanced

The question. Contact tracers find a new case. Should they look *forward*, for the people this case may have infected, or *backward*, for whoever infected them?

```
makeFromCSV("outbreak-transmissions-SIMULATED.csv",
 ↪  "outbreak-cases-SIMULATED.csv", directed=true)
o = offspring(show=False)
import statistics
round(o.R, 2), round(statistics.mean(e.node1.offspring for e in g.edges), 1)
```

A typical case infected 0.88 people (R over the whole outbreak). The person who infected a typical case had infected 5.6. That's the friendship paradox once more: a case is most likely to come from a large cluster, so its infector is likely a superspreader, and finding the infector finds their other cases too. Forward tracing from a new case finds 0.88 more cases on average; backward tracing finds 5.6. When dispersion is high (small k), backward tracing finds many more cases for the same effort, as Endo et al. (2020) and Kojaku et al. (2021) argued for COVID-19.

Capstone

Capstone: an outbreak control room Advanced

The problem. A school health officer must choose a plan before the winter: vaccinate some pupils (whom?), keep classes apart, or both. This recipe builds a small control room:

- the school drawn by class, each person coloured by their risk of infection under the current plan;
- a drop-down for the strategy and sliders for vaccine coverage and how contagious the infection is;

- a button that runs 200 outbreaks and reports the outcome;
- click a person to watch one outbreak start from them;
- a log of each plan tried, to compare them.

```
makeFromGDF("school-contacts-SIMULATED.gdf")
import random
ss("school")
plan = {"immune": [], "log": []}

def choose(strategy, share):
    dose = round(share * len(g.nodes))
    rng = random.Random(1)
    if strategy == "random":
        return rng.sample(list(g.nodes), dose)
    if strategy == "most contacts":
        return sortBy(degree)[len(g.nodes) - dose:]
    if strategy == "acquaintance":
        acq = set()
        while len(acq) < dose:
            acq.add(rng.choice(list(rng.choice(list(g.nodes)).getNeighbors())))
        return list(acq)
    return []

def setup():
    ls("school")
    if bubbles.checked:
        remove([e for e in g.edges if e.node1.room != e.node2.room])
    plan["immune"] = choose(strategy.value, coverage.value / 100)

def run():
    setup()
    beta = contagion.value / 100
    sz = outbreaks(200, beta=beta, gamma=0.2, immune=plan["immune"] or None)
    colorBy(risk, white, red)
    for n in plan["immune"]: n.color = "60,120,220"
    line = (f"{strategy.value}, {coverage.value:.0f}% vaccinated" + (", bubbles"
↪   if bubbles.checked else "")
            + f", beta {beta:.2f}: mean {sz.mean:.0f} infected, {sz.large:.0%}
             ↪  take off")
    plan["log"].append(line)
    result.text = "\n".join(plan["log"][-5:])
    StatusBar.setStatus(line)

def one_outbreak(n):
    if n in plan["immune"]:
        return
    spread([n], beta=contagion.value / 100, gamma=0.2, immune=plan["immune"] or
↪   None, seed=len(plan["log"]))
    colorBy(infected_at, yellow, red)
```

```
    (infected_at < 0).color = lightgray
    for m in plan["immune"]: m.color = "60,120,220"

strategy = Dropdown(["none", "random", "acquaintance", "most contacts"],
 ↪  value="acquaintance", label="vaccinate")
coverage = Slider(0, 50, 20, label="coverage %")
contagion = Slider(1, 10, 3, label="beta x 100")
bubbles = Checkbox("class bubbles", False)
result = Label("choose a plan and press Run")
ui.dock(Panel("Outbreak control room", strategy, coverage, contagion, bubbles,
              Button("Run 200 outbreaks", onClick=run), result,
 ↪  position="bottom"))
graphevents.clickNode = one_outbreak
run()
center()
```

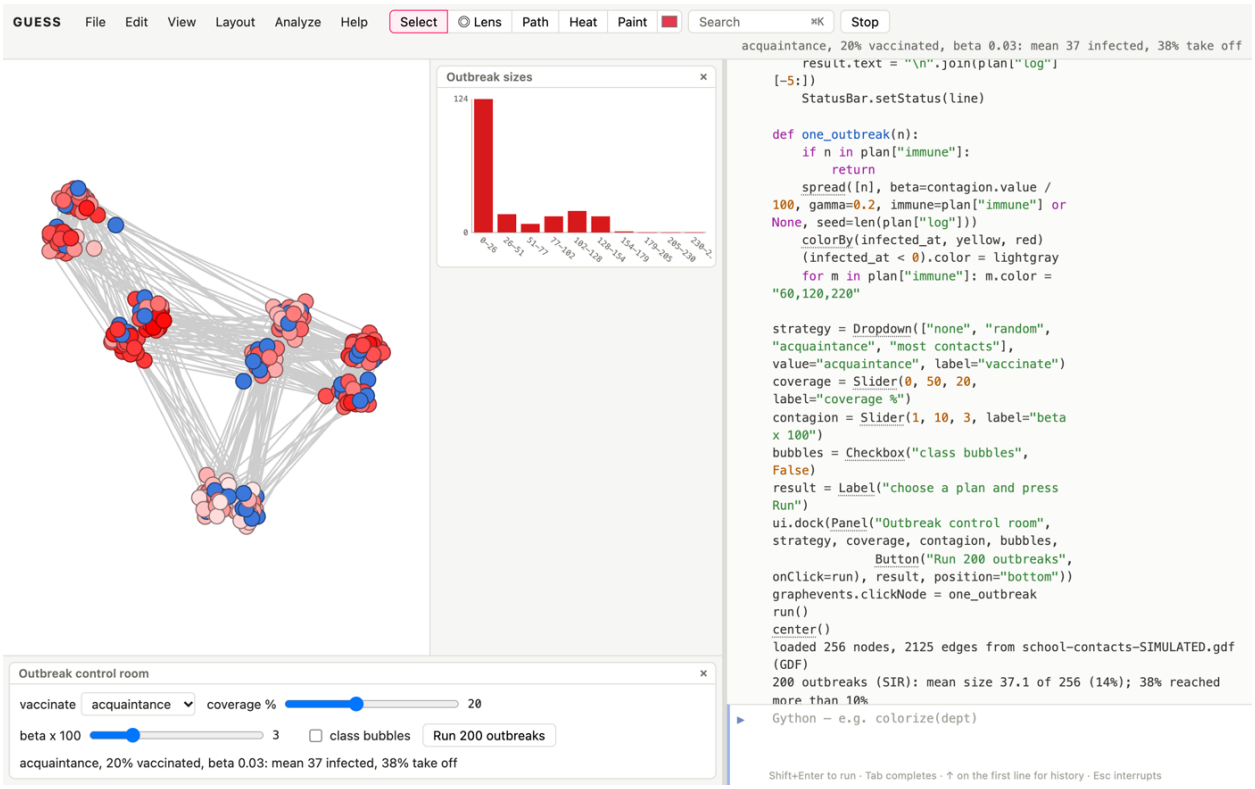


Figure 5: The control room: risk under acquaintance vaccination (vaccinated people in blue)

How it works.

- **The plan.** `setup()` restores the school with `ls`, cuts the between-class contacts if *class bubbles* is ticked, and chooses whom to vaccinate with the same strategies as the recipe above.
- **The button.** `run()` calls `outbreaks` with the vaccinated people as `immune`, colours everyone by `risk` (vaccinated people blue), and adds a line to the log. The outbreak-size chart updates in place.

- **Clicking.** `graphevents.clickNode = one_outbreak` starts a single outbreak from the clicked person under the current plan, coloured by day of infection, with its epidemic curve in the chart panel.

Variations.

- Add a slider for vaccine efficacy: drop a random share of `plan["immune"]` before each run.
- Add “teachers only” as a strategy (`role == "teacher"`).
- `exportApp()` shares the control room as a page.